



Hierarchical Reinforcement Learning (HRL): An Introduction

Barry Plunkett (eplu@sas.upenn.edu)

April 4, 2019

Problem & Motivation – Weaknesses of RL

- Sample inefficiency
- Sparse reward environments
- Large state, action space environments
- Unintuitive
- Generalization and abstraction
- Hold on...we solved Go!



Problem & Motivation – Promise of HRL

Hierarchical
Reinforcement
Learning

=

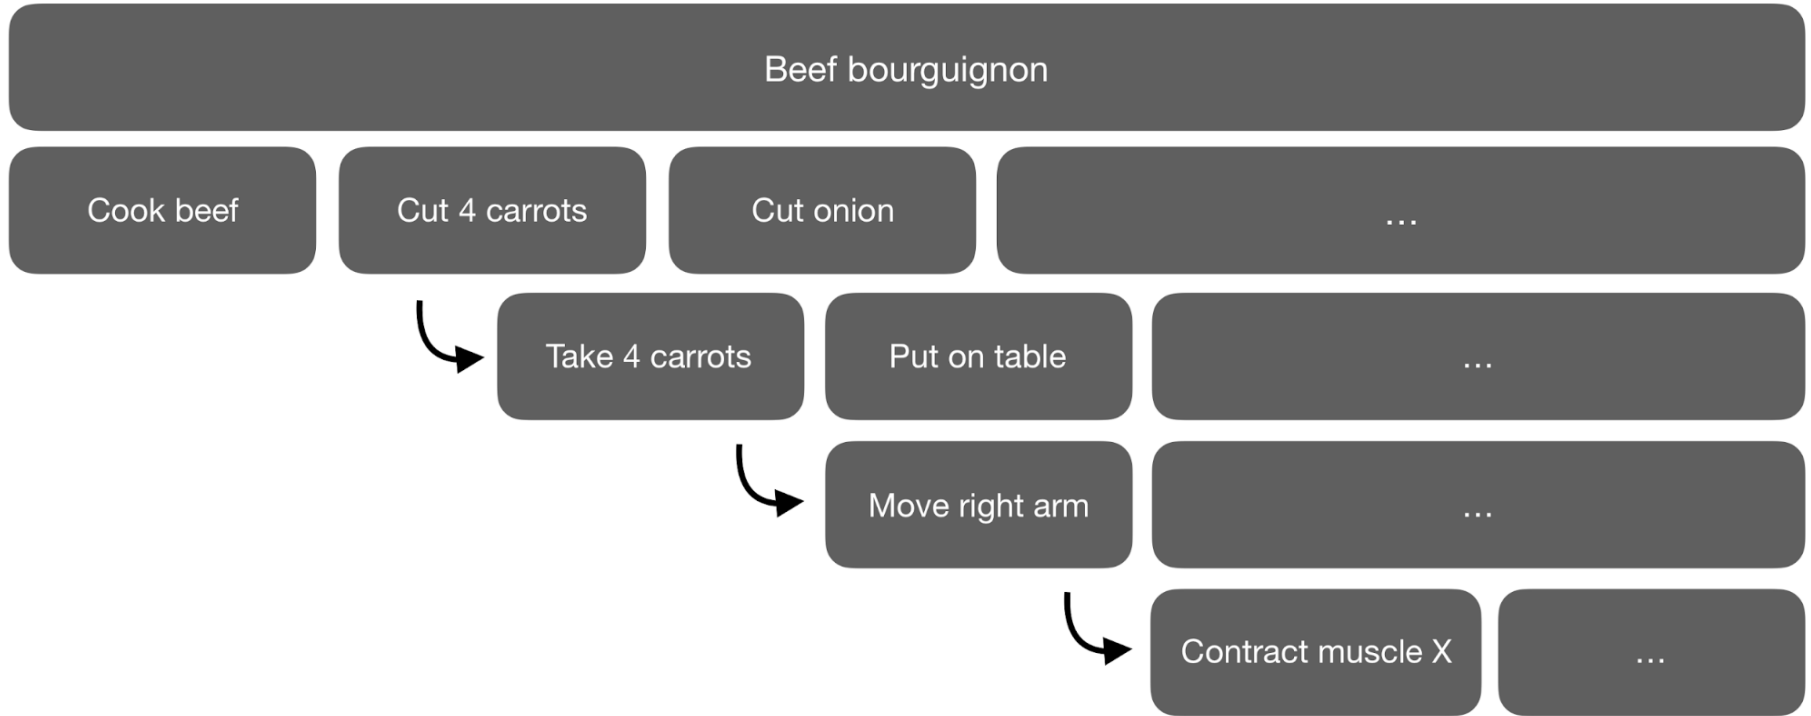
Reinforcement
Learning

+

Temporal
Abstraction

- Decompose goal into subtasks
 - Learn low-level policies to solve small subtasks
 - Compose low-level policies into longer-term, more abstract strategies to achieve goal
- Denser rewards
- Transfer learning via subproblem re-use
- Distill state/action space into cohesive subspaces

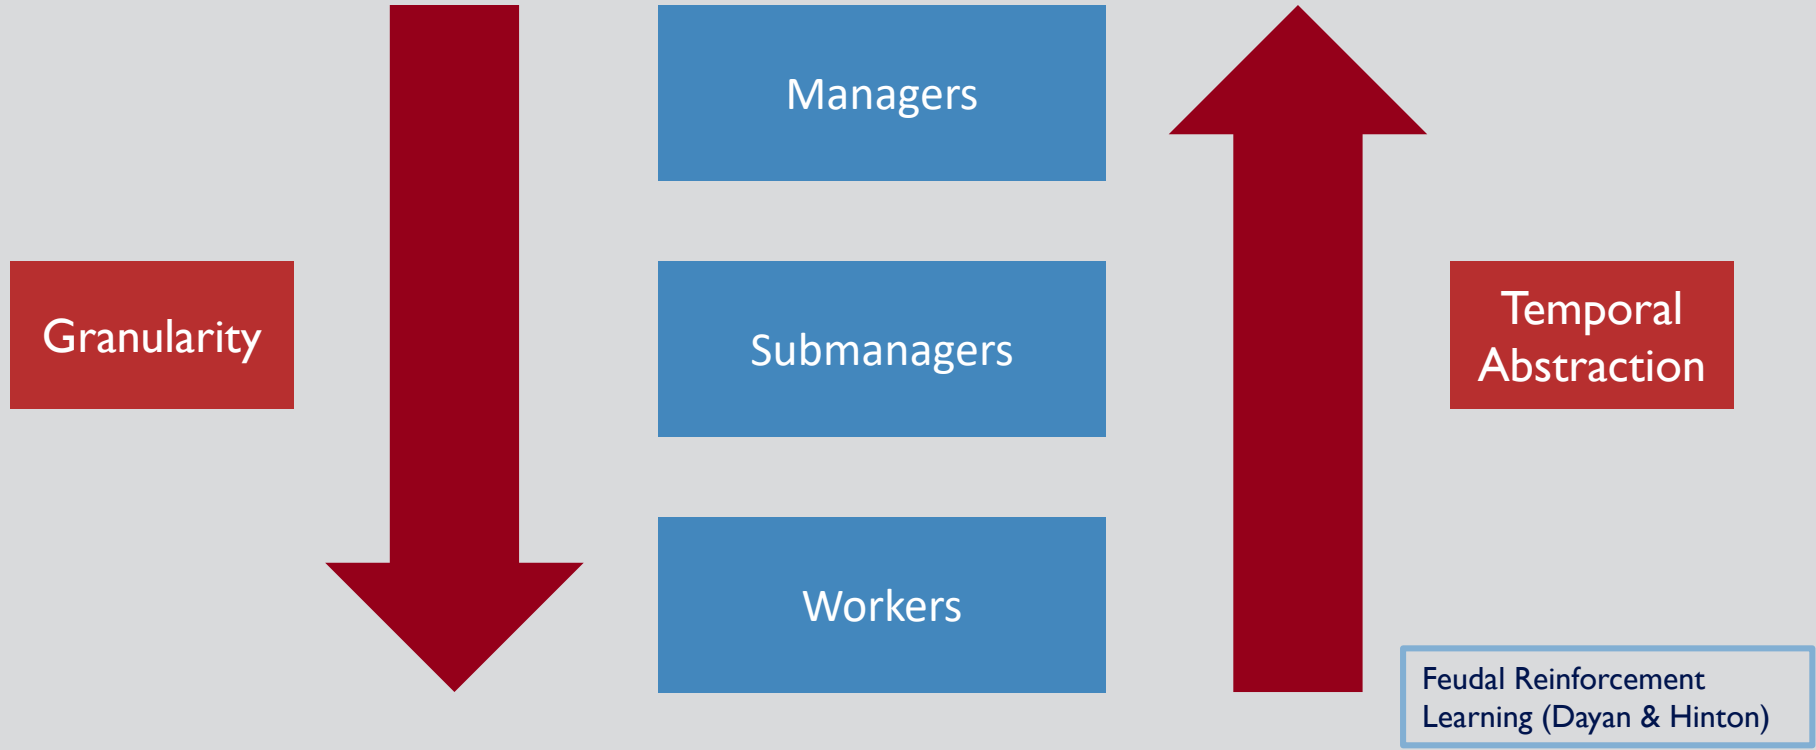
Problem & Motivation – Promise of HRL



Contents:

- Feudal Learning
- Markov Options
- HAMs
- MAXQ
- Conclusions

Feudal Learning – Introduction



Feudal Learning – Key Features


Reward Hiding

- A submanager receives reward if and only if it achieves the goal set for it by its manager
- No reward if manager goal attained but not submanager goal attained
- Receives reward if goal attained but manager goal not attained

Information Hiding


- Information hidden downwards – submanagers do not know their manager's task
- Information hidden upwards – managers do not know how their sub-managers have assigned workers to complete task

Feudal Learning – Maze Task



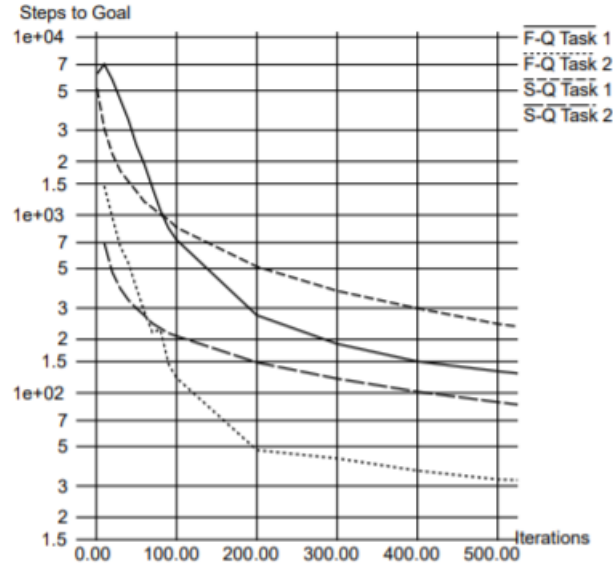
- N, S, E, W: Move to region in given cardinal direction at current level
- *: pass control to sub-managers to search for goal within current region at finer grain
- $A_1: \{*\}$
- $A_{1-(n-1)}: \{N, S, E, W, *\}$
- $A_n: \{N, S, E, W\}$
- State:
 - Action selected by manager above
 - Location of agent on the board in the granularity below

Feudal Learning – Maze Task



Tabular Q-values updated at all levels where a transition occurred, if the transition was ordered at all lower levels

Feudal Learning – Maze Task



- F-Q: Feudal system
- S-Q: Standard tabular Q
- Feudal slower initially
- Faster later

Feudal Learning – Conclusions


Advantages

- Learns more about environment than standard Q-Learning approach
- Structured exploration

Costs

- Information hiding may introduce inefficiencies
- Submanagers learn solutions to subproblems, even if these are not relevant to goal
- Task may appear non-markovian at high level abstraction

Markov Options – Introduction



Markov Options – Semi-Markov Decision Processes (SMDP)

- MDP : Amount of time between decisions fixed
- SMDP: Amount of time between decisions is random variable (τ)
 - Continuous
 - Discrete
- Treat system as “waiting” for τ periods
- Instantaneous state transition afterward



Markov Options – Semi-Markov Decision Processes (SMDP)

MDP

$$P(s'|s, a)$$

$$r(s, a)$$

$$V^*(s) = \max_a [R(s, a) + \gamma (\sum_{s'} P(s'|a, s) V^*(s'))]$$

$$Q^*(s, a) = R(s, a) + \gamma \sum_{s'} P(s'|a, s) \max_{a'} Q^*(s', a')$$

SMDP

$$P(s', \tau | s, a)$$

$$R(s, a)$$

$$V^*(s) = \max_a [R(s, a) + (\sum_{s', \tau} \gamma^\tau P(s', \tau | a, s) V^*(s'))]$$

$$Q^*(s, a) = R(s, a) + \sum_{s', \tau} \gamma^\tau P(s', \tau | a, s) \max_{a'} Q^*(s', a')$$

Markov Options - SMDP – Q-Learning

- Control via V, Q-learning again

$$Q^\pi(s,a) = Q^\pi(s,a) + \alpha(r_{t+1} + \gamma r_{t+2} + \dots + \gamma^\tau r_{t+\tau} + \max_{a'} Q^\pi(s', a') - Q^\pi(s,a))$$

- Converges under same guarantees as MDP Q-Learning
 - Linear function approximator
- Symmetric update for new value function too

Markov Options Formalization

Option (O)		
Input set	Policy	Termination Condition
$I \subseteq S$ Set of states where option O is available	$\pi: (S, A) \rightarrow [0, 1]$ Distribution of actions taken by options	$\beta: (S) \rightarrow [0, 1]$ Probability option ends in a given state

Markov Options – Assumptions

- Actions of core MDP
 - “Primitive actions” or one-step options

$$\beta(s) = 1 \quad \forall s \in S$$

- At least one option available in all states
- Option available in all states where it may continue

$$\{s : \beta(s) < 1\} \subseteq I$$



Open-the-door

Markov Options – Semi-Markov Options

- Semi-Markov options: Options where actions may depend on entire history of observations, since beginning of option

$$\mu : S \times \cup_{s \in S} O_s \rightarrow [0, 1]$$

- Allow options that terminate after fixed number of steps
- Allow policies over options
- Flat policies
 - Policy over primitive actions of core MDP
 - All μ correspond to some flat policy $flat(\mu)$
 - Non-Markovian even when all policies are Markovian

Markov Options – SMDP Q-learning

Reward	$R(s, o) = E[r_{t+1} + \gamma r_{t+2} + \dots + \gamma^{\tau} r_{t+\tau}]$
Transition Function	$P(s' s, o) = \sum_{t=1}^{\tau} p(s', t) \gamma^t$

$$V_O^*(s) = \max_{o \in O_s} \left[R(s, o) + \sum_{s'} P(s'|s, o) V_O^*(s') \right] \quad Q_O^*(s, o) = R(s, o) + \sum_{s'} P(s'|s, o) \max_{o' \in O_{s'}} Q_O^*(s', o')$$

$$Q_{k+1}(s, o) = (1 - \alpha_k) Q_k(s, o) + \alpha_k \left[r + \gamma^{\tau} \max_{o' \in O_{s'}} Q_k(s', o') \right]$$

Markov Options – Intra-option learning

- Q-Learning Drawbacks
 - Updates 1 option at a time
 - Must wait until option completes to update
- Intra-option learning methods
 - Learn online while option executes
- One-step intra-option Q-Learning
 - Suppose primitive action a taken, then for every option whose policy could have selected a with the same distribution $\pi(s, *)$:

$$Q_{k+1}(s_t, o) = (1 - \alpha_k)Q_k(s_t, o) + \alpha_k[r_{t+1} + \gamma U_k(s_t, o)].$$

$$U_k(s, o) = (1 - \beta(s))Q_k(s, o) + \beta(s) \max_{o' \in O} Q_k(s, o').$$

Markov Options – Conclusions

- Add temporally-extended activities without precluding fine-grained control
- Exclude some primitives
 - Restrict set of learnable policies
 - Increase efficiency, prevent “flailing”
- Utilize options to achieve subgoals
 - Define subgoal-specific reward functions and use for option policy
 - Set options to terminate upon subgoal completion



Hierarchies of Abstract Machines (HAMs)

- Apply temporal abstraction to **SIMPLIFY** rather than augment
- Well-known set of optimal (or good enough) policies for long-time horizon actions
 - Robot navigation




HAMs – Formalization

- Collection of finite state machines $\{H_i\}$
- Core environment MDP (M)
- Machine state initialization function $I_i: S_M \rightarrow S$
- Stochastic state transition function: $\delta_i S \rightarrow S$

State set (S)			
Action	Call	Choice	Stop
Executes action of M $a = \pi(m_t, s_t)$	Suspends execution of current machine, calls another machine H_j , function of m_t $I_j(s_t)$ sets initial state	Non-deterministically chooses next state of H_i	Suspends execution of current machine. Resumes execution of calling machine

HAMs

- If no action is generated at step t , M remains in current state
- H : Initial machine
 - Assume no stop state
 - Assume no probability 1 loops
 - Ensure MDP will continue to receive primitive actions



HAMs – SMDP view

- Equivalent to SMPD : $H \circ M$
- State : $S \times S_M$
- Actions: Choice points of H
 - Runs autonomously via action states until next choice point reached
 - Do not learn within machine policies – program these
- Reward: Discounted awards accumulated during timesteps between choice states
 - Reward of 0 for timesteps where M does not change

HAMs – Q-learning

- $Reduce(H \circ M)$
 - SMDP equivalent to $H \circ M$ with states defined as only choice points of $H \circ M$
 - Optimal policy for $Reduce(H \circ M)$ same as $H \circ M$
- Apply standard SMDP q-learning update to $Reduce(H \circ M)$


$$Q_{k+1}([s_c, m_c], a_c) = (1 - \alpha_k)Q_k([s_c, m_c], a_c) + \alpha_k[r_{t+1} + \gamma r_{t+2} + \cdots + \gamma^{\tau-1}r_{t+\tau} + \gamma^{\tau} \max_{a'} Q_k([s'_c, m'_c], a')]$$

HAMs - Conclusions

- Success depends on quality of programmed policies and state transition functions for each machine
- Not used in any large-scale applications
- Allows integration of multiple controls problems, whose solutions are well known, but whose relationships are not

MAXQ

Decompose into hierarchy of SMDPs (rather than 1) and solve simultaneously



MAXQ – Formalization

- Decompose MDP M into subtasks $M_0 \dots M_n$
 - M_0 corresponds to original task
- Each subtask similar to an option

Subtask anatomy (M_i)		
Policy	Active States	Pseudo-Reward
$a = \pi(s_t, k)$ <ul style="list-style-type: none">• Assume deterministic• K denotes subtask call stack	<ul style="list-style-type: none">• S_i: Set of states where M_i can execute• T_i: ($S \setminus S_i$) states where subtask terminates	<ul style="list-style-type: none">• Task-specific reward function that assigns reward to each state in T_i

MAXQ – SMPDs

- Representation of top-level policy

$$\pi = \{\pi_0, \dots, \pi_n\}$$

- Transition probabilities

$$P_i(s', \tau | s, a)$$

- Value of completing i th subtask task from state s and following π

$$V^\pi(i, s)$$

MAXQ – State Value Function

- Reward of selecting subtask a from subtask i :

$$R_i(s, a) = V^\pi(a, s)$$

- Corresponding Bellman state value equation

$$V^\pi(i, s) = V^\pi(\pi_i(s), s) + \sum_{s', \tau} P_i^\pi(s', \tau | s, \pi_i(s)) \gamma^\tau V^\pi(i, s')$$

MAXQ – Action-Task Value Function

Q-Function

$$Q^\pi(i, s, a) = V^\pi(a, s) + \sum_{s', \tau} P_i^\pi(s', \tau | s, a) \gamma^\tau Q^\pi(i, s', \pi(s'))$$

$$C^\pi(i, s, a) = \sum_{s', \tau} P_i^\pi(s', \tau | s, a) \gamma^\tau Q^\pi(i, s', \pi(s'))$$


$$Q^\pi(i, s, a) = V^\pi(a, s) + C^\pi(i, s, a)$$

$$V^\pi(0, s) = V^\pi(a_n, s) + C^\pi(a_{n-1}, s, a_n) + \cdots + C^\pi(a_1, s, a_2) + C^\pi(0, s, a_1)$$

$$V^\pi(a_n, s) = \sum_{s'} P(s' | s, a_n) R(s' | s, a_n)$$

MAXQ Learning

- High-level overview of complex algorithm
 - Similar to Monte Carlo Q-learning with completion function
- Estimate completion $C(i, s, a)$ for each $i \rightarrow a$ edge in the tree
- Beginning at 0, recursively execute actions to descend to primitive MDP (choosing subtask with highest completion estimate)
- After each subtask a concludes, use reward accumulated at leaf below to update $C(i, s, a)$



MAXQ Conclusions

- Recursively optimal policy
 - Optimal for a given subtask SMDP, given SMDPs of children
- Weaker than hierarchically optimal policy
 - Optimal policy among all policies that can be expressed within constraints of hierarchy
- Why? Hierarchically optimal policy may need to exploit context of calling subtask
 - E.g. Optimal way to travel to a destination may depend on what you're doing at the destination



Conclusions

- HRL allows temporal abstraction for long-term and short-term planning –maps onto human decision making
- Potentially valuable avenue for improving transfer learning of tasks and strategies
- Unsolvable problems today have characteristics HRL seems well suited to address
 - Sparse rewards
 - Large action and state spaces
 - Slow data generation

Conclusions

- Fundamental Problem: All frameworks we've seen today require hand-generated decompositions
 - Deal-breaker for complex tasks
 - ...and that's the whole point
- Learning decompositions automatically is an active area of research
 - Heuristics & approximations
 - MetaRL (approaches reminiscent of AutoML)
- Marriages with deep RL increasingly common
 - FUN : Deep feudal learning
 - Option-Critic: Actor-critic policy gradient setup meets options framework
 - HIRO, h-DQN, and many more